

1 Staged Pre-Training with Frozen Backbone and 2 Parameter-Efficient Fine-Tuning for Financial RL Portfolio 3 Optimization 4

5 Anonymous Author(s)
6
7

8 Abstract

9 Reinforcement learning (RL) has shown promise in portfolio management, particularly in handling customizable stock pools (CSPs).
10 However, prevailing joint end-to-end training approaches, exemplified by EarnMore, suffer from severe gradient interference in financial environments characterized by extremely sparse, noisy, and delayed rewards. This often leads to unstable representation learning, poor generalization, and degraded performance under market regime shifts. To address these challenges, we propose SPF (Staged Pre-training with Frozen backbone), a decoupled, staged framework tailored for low-signal financial RL. First, we pre-train a robust, mask-aware representation via multi-task self-supervision, combining masked reconstruction with latent-space contrastive learning to capture both local price patterns and inter-stock semantic relationships. The pre-trained backbone is then fully frozen to preserve financial inductive biases and prevent destructive interference from weak RL gradients. Finally, we perform highly parameter-efficient fine-tuning using a hierarchical residual lightweight adapter chain (inspired by Lossless Adaptation) at multiple levels of abstraction, while confining RL updates to only the adapters, SAC critic, and entropy coefficient (with optional actor adaptation). Extensive experiments on the DJ30 and SP500 datasets, including Global Stock Pool (GSP) and multiple CSP benchmarks, demonstrate that SPF significantly outperforms joint-training baselines, achieving substantial improvements in risk-adjusted metrics (e.g., Sharpe ratio, cumulative return, Calmar ratio) and enhanced robustness across out-of-sample periods and market regimes. The proposed approach also exhibits lower variance and better adaptation to transaction costs. Our code is publicly available at <https://github.com/Flying932/PFRL>.

39 CCS Concepts

- 40 • Computing methodologies → Reinforcement learning; Neural networks; Machine learning; • Applied computing → Economics.

44 Keywords

46 Reinforcement Learning, Portfolio Management, Customizable Stock
47 Pools, Self-Supervised Pre-training, Frozen Representations, Parameter-
48 Efficient Fine-Tuning, Financial Time Series

49
50 Permission to make digital or hard copies of all or part of this work for personal or
51 educational use is granted. To copy otherwise, or to redistribute in whole or in part,
52 without prior permission or written license, to do so must be distributed for profit or
53 commercial advantage and that copies bear this notice and the full citation
54 on the first page. Copyrights for components of this work owned by others than the
55 author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or
56 republish, to post on servers or to redistribute to lists, requires prior specific permission
57 and/or a fee. Request permissions from permissions@acm.org.

58 ACM MM '26, Rio de Janeiro, Brazil

59 © 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
60 ACM ISBN 978-1-4503-XXXX-X/18/06
61 <https://doi.org/XXXXXXX.XXXXXXXX>

62 2026-02-10 13:39. Page 1 of 1–5.

63 ACM Reference Format:

64 Anonymous Author(s). 2026. Staged Pre-Training with Frozen Backbone and
65 Parameter-Efficient Fine-Tuning for Financial RL Portfolio Optimization.
66 In *Proceedings of Proceedings of the 34th ACM International Conference on
67 Multimedia (ACM MM '26)*. ACM, New York, NY, USA, 5 pages. <https://doi.org/XXXXXXX.XXXXXXXX>

68 1 Introduction

69 Portfolio management under customizable stock pools (CSPs) constitutes a challenging sequential decision-making task in which the
70 available asset universe changes dynamically [cite: 58, 79]. Recent
71 reinforcement learning (RL) methods, particularly EarnMore [?], have advanced this domain by introducing maskable representations
72 that enable a single model to adapt to varying stock pools [cite: 60, 80]. Despite this progress, such approaches rely predominantly
73 on joint end-to-end training of representation learning and policy optimization [cite: 59, 81].

74 Financial RL environments exhibit extreme characteristics: rewards are sparse, noisy, and significantly delayed, while the underlying
75 data distribution is highly non-stationary [cite: 59, 82, 114]. A fundamental challenge arises from the deep misalignment
76 between the dense signals in self-supervised learning and the sparse, high-variance rewards in RL. While pre-training
77 objectives like price reconstruction provide a dense supervisory signal across all tokens—enabling the model to capture the intricate
78 "market physics"—RL rewards are often dominated by idiosyncratic market noise. Under joint training, early excessive exploration
79 involves stochastic actions that generate high-variance gradients; these gradients, when backpropagated, can effectively
80 "reset" the encoder's weights. Consequently, the model loses its painstakingly learned sensitivity to critical price trends and
81 inter-stock relationships, leading to unstable representations and substantial performance degradation during market regime shifts.

82 To address these issues, we adopt a staged pre-training and fine-tuning paradigm that deliberately decouples representation learning
83 from policy optimization [cite: 60, 86, 370]. By treating the pre-training as a dense knowledge-acquisition stage, we establish
84 a robust and stable initialization that captures temporal dynamics and market patterns from unlabeled data [cite: 61, 87].
85 This stage remains largely unaffected by the sparsity and noise of subsequent RL rewards [cite: 88, 119]. During fine-tuning, the
86 pre-trained Transformer backbone is fully frozen to protect these high-quality representations from the destructive interference of
87 sparse RL gradients [cite: 62, 89, 192]. This "freeze-and-adapt" strategy ensures that the RL agent builds upon a fixed, reliable
88 financial foundation rather than corrupting it through noisy updates.

89 Adaptation to the downstream RL task occurs through a hierarchical residual lightweight adapter chain (inspired by Lossless
90 Adaptation) at multiple levels of abstraction, while confining RL updates to only the adapters, SAC critic, and entropy coefficient
91 (with optional actor adaptation). Extensive experiments on the DJ30 and SP500 datasets, including Global Stock Pool (GSP) and
92 multiple CSP benchmarks, demonstrate that SPF significantly outperforms joint-training baselines, achieving substantial
93 improvements in risk-adjusted metrics (e.g., Sharpe ratio, cumulative return, Calmar ratio) and enhanced robustness across
94 out-of-sample periods and market regimes. The proposed approach also exhibits lower variance and better adaptation to
95 transaction costs. Our code is publicly available at <https://github.com/Flying932/PFRL>.

117 Adaptation [?]), inserted at multiple levels of abstraction[cite: 63,
 118 193, 451]. RL updates are strictly confined to the adapters, SAC critic
 119 networks, and entropy coefficient, ensuring minimal disruption to
 120 the frozen financial priors[cite: 63, 194, 619].

121 The remainder of this paper is organized as follows: Section 2
 122 reviews related work, Section 3 presents the proposed framework,
 123 Section 4 reports experimental results, and Section 5 concludes the
 124 paper.
 125

2 Related Work

2.1 Joint Representation Learning with RL

126 Joint optimization of state representations and RL policies has been
 127 a dominant approach in recent portfolio management research.
 128 Yang et al. (2023) highlighted that many existing RL methods treat
 129 assets in isolation and optimize solely for episodic rewards, neglecting
 130 the global portfolio context. They introduced the Task-Context
 131 Mutual Actor-Critic (TC-MAC) algorithm, which simultaneously
 132 learns a task-context representation encoding both individual asset
 133 features and the overall dynamic portfolio context, along with a
 134 policy optimized through a combined RL and mutual-information
 135 loss. By maximizing mutual information between local asset em-
 136 beddings and the global context, TC-MAC achieves more coherent
 137 portfolio representations and improved policy optimality, outper-
 138 forming traditional methods and prior RL baselines with enhanced
 139 transferability to unseen market data.
 140

141 EarnMore [?] represents a key advancement in this direction. To
 142 address the limitation of fixed asset universes, it performs one-shot
 143 training on a large global stock pool and proposes maskable stock
 144 representations that adapt to arbitrary customizable subsets. The
 145 method employs a self-supervised autoencoder for masked recon-
 146 struction and uses learnable mask tokens to suppress irrelevant
 147 stocks, combined with a re-weighting mechanism to concentrate
 148 the portfolio on promising assets. The authors adopt a three-stage
 149 joint end-to-end training pipeline and argue that freezing pretrained
 150 embeddings may restrict RL exploration and task alignment. Earn-
 151 More achieved substantial profit gains (over 40%) over 14 baseline
 152 methods on U.S. stock datasets. However, joint training exposes
 153 the representation module to weak and noisy RL gradients, poten-
 154 tially leading to sample inefficiency and instability in reward-sparse
 155 financial environments.
 156

2.2 Incorporating Financial Knowledge and Hybrid Training

157 Several works integrate domain knowledge or supervised signals to
 158 improve sample efficiency and alignment with financial objectives.
 159 Hu and Gu (2024) combined Modern Portfolio Theory with deep RL
 160 in a two-stage framework. The first stage distills the agent on signals
 161 from Markowitz optimal portfolios (balancing return and risk),
 162 while the second stage performs RL fine-tuning. This knowledge-
 163 distillation approach (KDD) provides a strong warm start and yields
 164 higher returns and risk-adjusted performance (Sharpe ratio of 2.03)
 165 compared to pure RL and traditional baselines.
 166

167 Li et al. (2025) proposed Logic-Q, a neural-symbolic framework
 168 that augments deep RL with rule-based trend detection for quanti-
 169 tative trading. By injecting lightweight program sketches encoding
 170

171 human-defined logic (e.g., macro trend indicators), Logic-Q en-
 172 ables the agent to adjust risk exposure or actions based on market
 173 regimes. Evaluations demonstrated significant reductions in cata-
 174 strophic drawdowns during downturns and improved exploitation
 175 in uptrends, outperforming prior DRL strategies.
 176

2.3 Representation Learning for Stability and Generalization

177 Representation learning has become central to enhancing RL stabil-
 178 ity and generalization in financial environments. Wu et al. (2023)
 179 addressed overfitting and volatility by combining contrastive learn-
 180 ing with reward smoothing. The contrastive objective encourages
 181 trend-aware state representations that cluster similar long-term
 182 dynamics and separate dissimilar ones, while reward smoothing
 183 mitigates short-term noise. This approach improved consistency
 184 and robustness on both stock and cryptocurrency markets.
 185

186 Other efforts leverage graph neural networks to capture cross-
 187 asset correlations and sector dynamics, or attention mechanisms
 188 to dynamically focus on relevant assets and market signals. These
 189 representation-oriented techniques consistently demonstrate im-
 190 proved robustness and generalization compared to raw price-based
 191 state representations.
 192

193 In summary, recent advances have significantly improved per-
 194 formance through richer representations, domain knowledge inte-
 195 gration, auxiliary tasks, and hybrid training schemes. Nevertheless,
 196 most methods rely on joint end-to-end optimization or partial pre-
 197 training, leaving room for fully decoupled paradigms that protect
 198 high-quality representations from sparse RL interference while en-
 199 abling efficient task-specific adaptation. Our work builds on these
 200 insights by proposing a staged, frozen-backbone approach with
 201 multi-task self-supervision and hierarchical adapters, specifically
 202 tailored to the low-signal, non-stationary nature of financial port-
 203 folio optimization.
 204

3 Method

205 We propose SPF (Staged Pre-training with Frozen backbone), a
 206 decoupled framework that separates dense self-supervised repre-
 207 sentation learning from sparse RL optimization for robust portfolio
 208 management in customizable stock pools (CSPs). The framework
 209 consists of three main stages: (1) multi-task self-supervised pre-
 210 training, (2) complete backbone freezing, and (3) parameter-efficient
 211 fine-tuning with RL adaptation.
 212

3.1 Stage 1: Multi-Task Self-Supervised Pre-Training

213 The goal of the pre-training stage is to learn high-quality, mask-
 214 aware stock representations that capture both local temporal pat-
 215 terns and inter-stock semantic relationships, providing a strong
 216 and stable initialization for downstream RL.
 217

218 We build directly on the self-supervised tasks introduced in
 219 EarnMore [?]. Let $X \in \mathbb{R}^{B \times T \times N \times D}$ denote a batch of stock time-
 220 series sequences, where B is the batch size, T is the look-back
 221 window, N is the number of stocks in the global pool, and D is the
 222 feature dimension (OHLCV + technical indicators).
 223

224 225 226 227 228 229 230 231

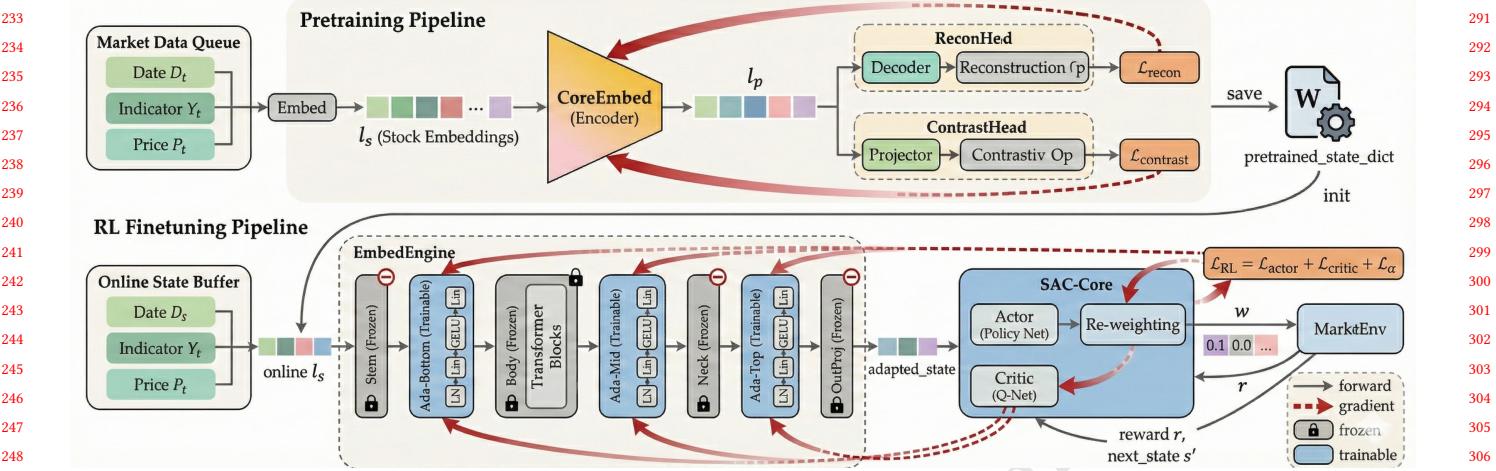


Figure 1: Overview of the SPF (Staged Pre-training with Frozen backbone) framework. Stage I (top) illustrates the multi-task self-supervised pre-training pipeline, where the encoder (CoreEmbed) learns from both reconstruction and contrastive signals. Stage II (bottom) shows the RL fine-tuning process. The pre-trained backbone is frozen to preserve financial semantic knowledge, while hierarchical residual adapters (Ada-Bottom, Ada-Mid, and Ada-Top) are optimized alongside the SAC agent to adapt to specific portfolio management tasks in customizable stock pools.

3.1.1 Masked Reconstruction (Inherited from EarnMore). We randomly mask a portion (mask ratio $\rho \in [0.15, 0.3]$) of the input tokens along the time and stock dimensions. Let M be the binary mask matrix with the same shape as X , where $M_{b,t,n,d} = 1$ indicates a masked position. The encoder $E(\cdot)$ (a Transformer-based architecture) processes the masked input $X \odot (1 - M) + m \cdot M$, where m is a learnable mask token. The decoder $D(\cdot)$ reconstructs the original values at masked positions:

$$\hat{X} = D(E(X \odot (1 - M) + m \cdot M)) \quad (1)$$

The reconstruction loss is the mean squared error over masked positions only:

$$\mathcal{L}_{\text{recon}} = \frac{1}{\sum M} \sum_{b,t,n,d} M_{b,t,n,d} \cdot (\hat{X}_{b,t,n,d} - X_{b,t,n,d})^2 \quad (2)$$

This task encourages the encoder to learn local temporal dynamics and stock-specific patterns.

3.1.2 Latent Contrastive Learning (Our Extension). To enrich the representations with robust semantic consistency, we introduce a dual-view contrastive learning objective in the latent space. Unlike traditional stock-pair contrast, our approach constructs views through stochastic perturbations. Specifically, for each mini-batch of market state windows $X \in \mathbb{R}^{B \times T \times N \times D}$, we perform two independent forward passes through the encoder $E(\cdot)$ with different random masking seeds (maintaining the same mask ratio ρ). This yields two distinct views of the same batch: $Z^{(1)} = E(X; \text{seed}_1)$ and $Z^{(2)} = E(X; \text{seed}_2)$, where $Z^{(k)} \in \mathbb{R}^{B \times C}$ represents the global latent embedding after temporal and pool-level pooling.

Positive and Negative Pairs. For each anchor embedding $z_i^{(1)}$ in $Z^{(1)}$, we define its corresponding embedding from the same market

state, $z_i^{(2)}$, as the unique positive pair. The remaining $2B - 2$ embeddings within the combined set $\{Z^{(1)}, Z^{(2)}\}$ (excluding the anchor and its positive pair) are treated as negative samples, representing different time windows or market regimes.

Contrastive Objective. Following the NT-Xent (Normalized Temperature-scaled Cross Entropy) framework, we apply L2-normalization to all embeddings and compute similarity via dot products[cite: 378, 382]. The loss for a single anchor i is formulated as:

$$\ell_i = -\log \frac{\exp(z_i^{(1)\top} z_i^{(2)}/\tau)}{\sum_{k=1}^B [\exp(z_i^{(1)\top} z_k^{(1)}/\tau) + \exp(z_i^{(1)\top} z_k^{(2)}/\tau)] - \exp(z_i^{(1)\top} z_i^{(1)}/\tau)} \quad (3)$$

where $\tau = 0.2$ is the temperature coefficient. The total contrastive loss $\mathcal{L}_{\text{contrast}}$ is the average over all $2B$ samples in the dual-view batch[cite: 412, 443].

This task encourages the model to maintain semantic invariance under masking noise—emphasizing that the "same market state" should have a consistent representation regardless of partial occlusion—while distinguishing between different market conditions.

Discussion. While effective for robust feature extraction, we acknowledge that this masking-based contrast primarily captures stochastic invariance. Current negative samples from the same batch may contain "false negatives" (e.g., temporally adjacent or semantically similar windows)[cite: 191]. Future extensions could incorporate stronger financial inductive biases, such as industry-consistent or correlation-aware positive sampling[cite: 61, 92].

The total pre-training loss is thus:

$$\mathcal{L}_{\text{pre}} = \lambda_1 \mathcal{L}_{\text{recon}} + \lambda_2 \mathcal{L}_{\text{contrast}} \quad (4)$$

where we set $\lambda_1 = \lambda_2 = 1.0$ [cite: 443, 444].

Table 1: Performance comparison on SP500 and DJ30 with Global Stock Pool. Results in red, yellow and green show the best, second best and third best results on each dataset.

2*Categories	2*Strategies	SP500				DJ30			
		ARR↑	SR↑	CR↑	SOR↑	ARR↑	SR↑	CR↑	SOR↑
3*Rule-based	Market	9.320	0.556	0.702	17.120	6.710	0.458	0.776	15.960
	BLSW	11.630	0.696	0.894	21.560	7.610	0.512	0.858	16.560
	CSMBoost	5.070	0.329	0.434	9.840	5.930	0.400	0.643	12.950
2*ML-based	XGBoost	10.690	0.377	0.575	13.650	10.340	0.343	0.599	14.220
	LightGBM	16.330	0.575	0.744	13.110	10.320	0.591	0.703	14.220
2*DL-based	ALSTM	43.560	1.157	1.044	22.550	15.930	1.186	0.732	8.890
	TCN	13.560	0.875	0.744	16.350	6.980	0.186	0.269	8.490
5*RL-based	PG	12.580	0.431	0.519	24.370	7.920	0.321	0.435	8.040
	PPO	15.130	0.538	0.743	15.470	22.900	0.689	1.465	8.350
	SAC	15.034	0.538	0.673	14.570	20.900	0.326	0.445	8.670
	EIIIE	21.240	0.756	0.970	25.230	27.192	0.786	1.109	27.320
	DeepTrader	50.320	1.162	0.949	35.230	27.420	0.909	1.593	27.230
	SPF (Ours)	xxx	xxx	xxx	xxx	xxx	xxx	xxx	xxx

3.2 Stage 2: Complete Backbone Freezing

After pre-training, the entire Transformer encoder $E(\cdot)$ is frozen (parameters set to require no gradient). This step protects the high-quality financial representations learned from dense self-supervised signals from being corrupted by the sparse, noisy gradients of RL during fine-tuning.

3.3 Stage 3: Parameter-Efficient Fine-Tuning with RL Adaptation

In the fine-tuning stage, the frozen encoder produces representations $Z = E(X)$ for each input sequence. To adapt these representations to the RL task (optimal portfolio allocation) with minimal interference, we insert a hierarchical residual lightweight adapter chain at multiple levels of abstraction.

The adapter chain consists of three sequential residual adapters (bottom, mid, top):

- **Bottom Adapter:** operates on token-level embeddings
- **Mid Adapter:** operates on pooled stock-level representations
- **Top Adapter:** operates on global sequence representation

Each adapter applies a residual transformation: it first normalizes the input z via **Layer Normalization**, then processes it through a **low-rank bottleneck linear layer** (with **GELU activation**), followed by a second linear layer and **dropout**. Finally, the output is added back to the original input z through a **residual connection** to ensure stable information flow. This process is compactly represented as $\text{Adapter}(z) = z + \text{Dropout}(\text{Linear}_2(\text{GELU}(\text{Linear}_1(\text{LayerNorm}(z)))))$, where both linear layers use a low-rank projection (bottleneck dimension $r \ll C$, e.g., $r = 64$). The residual connection ensures that the adapter only learns task-specific corrections without overwriting the pre-trained features.

The adapted representation is then fed into the SAC agent. RL updates are strictly confined to: - all parameters of the hierarchical

adapters - the SAC critic networks (dual Q-functions) - the entropy coefficient α (automatic temperature tuning) - the actor network (optional; our main experiments freeze the actor to further minimize interference)

This parameter-efficient fine-tuning (PEFT) strategy ensures stable convergence in the low-signal financial environment while adapting the pre-trained representations to the specific allocation task.

The overall objective in the fine-tuning stage is the standard SAC loss:

$$\mathcal{L}_{\text{SAC}} = \mathcal{L}_{\text{critic}} + \mathcal{L}_{\text{actor}} + \mathcal{L}_{\alpha} \quad (5)$$

where $\mathcal{L}_{\text{critic}}$ is the Bellman error on the dual Q-functions, $\mathcal{L}_{\text{actor}}$ maximizes $Q - \alpha \log \pi$, and \mathcal{L}_{α} enforces target entropy.

Through this three-stage design, SPF achieves robust, efficient, and generalizable RL-based portfolio management tailored to the unique challenges of financial domains.

4 Experimental Evaluation

4.1 Datasets

We conduct experiments on two widely-used financial benchmarks: DJ30 and SP500. The DJ30 dataset consists of the 30 constituent stocks of the Dow Jones Industrial Average, while SP500 includes the constituents of the S&P 500 index. Both datasets cover the period from September 26, 2007 to June 26, 2022, with input features comprising 102 dimensions. The prediction targets are ret1 (one-day return) and mov1 (movement label).

The data is split chronologically as follows:

- **Pre-training stage:** Training set: September 26, 2007 – July 22, 2019; Validation set: July 22, 2019 – June 26, 2022.
- **Fine-tuning stage:** Training set: September 26, 2007 – July 22, 2019; Validation set: July 22, 2019 – January 8, 2021; Test set: January 8, 2021 – June 26, 2022.

465 4.2 Experimental Settings

466 All experiments are conducted on an NVIDIA RTX 4090 GPU with
 467 24GB memory. To ensure reliability and reduce variance, we per-
 468 form 3 independent data splits and 5 random seeds for each setting,
 469 reporting the average and standard deviation.

470 Key hyperparameters are as follows: Pre-training uses a mask
 471 ratio $\rho = 0.25$, contrastive temperature $\tau = 0.2$, and reconstruction
 472 weight $\lambda_1 = 1.0$. Fine-tuning employs the SAC algorithm with a
 473 learning rate of 3×10^{-5} and a buffer size of 10^6 . Other settings
 474 follow the default configurations of ElegantRL and Qlib.

476 4.3 Comparison with Baselines

477 We compare SPF against a range of baselines including rule-based,
 478 ML-based, and RL-based methods. Performance comparison un-
 479 der the Global Stock Pool (GSP) setting is detailed in Table 1. For
 480 Customizable Stock Pools (CSP), results are presented in Table 2.

482
 483 **Table 2: Performance comparison on SP500 and DJ30 with**
 484 **Customizable Stock Pools. Underlined metrics indicate the**
 485 **best-performing index results.**

488 Pool	Strategies	SP500		DJ30	
		ARR↑	SR↑	ARR↑	SR↑
491 CSP1	SARL (Baseline)	34.333	0.820	24.141	0.638
	IMIT (Baseline)	20.973	0.860	20.741	0.928
	SPF (Ours)	<u>122.610</u>	<u>2.278</u>	<u>53.990</u>	<u>1.810</u>
494 CSP2	SARL (Baseline)	17.000	0.578	20.020	0.820
	DeepTrader (Baseline)	34.030	0.793	38.470	0.955
	SPF (Ours)	<u>110.110</u>	<u>2.279</u>	<u>43.400</u>	<u>1.549</u>
498 CSP3	SARL (Baseline)	18.090	0.760	10.910	0.480
	IMIT (Baseline)	61.320	1.489	16.840	0.601
	SPF (Ours)	<u>93.670</u>	<u>2.120</u>	<u>43.460</u>	<u>1.572</u>

501 4.4 Ablation Study

502 To verify the effectiveness of each component, we conduct an abla-
 503 tion study with variants: (1) w/o Pre-training, (2) w/o Contrastive,
 504 and (3) w/o Reconstruction. Table 3 reports the results, where red
 505 and green indicate performance changes.

509
 510 **Table 3: Ablation Study of SPF on SP500 and DJ30. Red, green**
 511 **and underline indicate improvement, performance decrease**
 512 **and best results, respectively.**

515 2*Model	SP500			DJ30		
	ARR↑	SR↑	MDD↓	ARR↑	SR↑	MDD↓
w/o Pre-training	xxx	xxx	xxx	xxx	xxx	xxx
w/o Contrastive	xxx	xxx	xxx	xxx	xxx	xxx
w/o Reconstruction	xxx	xxx	xxx	xxx	xxx	xxx
Full SPF (Ours)	xxx	xxx	xxx	xxx	xxx	xxx

523 5 Conclusion

524 We introduced a staged pre-training and frozen representation
 525 framework for robust RL in customizable portfolio management.
 526 Experiments show significant improvements over joint-training
 527 baselines in financial environments with weak signals. Future work
 528 includes multi-modal extensions and real-world deployment.

529 References

530
 531
 532
 533
 534
 535
 536
 537
 538
 539
 540
 541
 542
 543
 544
 545
 546
 547
 548
 549
 550
 551
 552
 553
 554
 555
 556
 557
 558
 559
 560
 561
 562
 563
 564
 565
 566
 567
 568
 569
 570
 571
 572
 573
 574
 575
 576
 577
 578
 579
 580